Oleksii Krasnoshtanov

oleksiikrasnoshtanov@gmail.com | +31 6 29746862 | Tilburg, Netherlands
linkedin.com/in/oleksii-krasnoshtanov | github.com/alex-krasnoshtanov

---

October 18, 2026

Genmab, Early Translational Data Science
Utrecht, Netherlands

Dear Hiring Manager,

I am applying for the internship in the Early Translational Data Science team on modelling historical control-arm data from preclinical tumour studies. I am a third-year student in Applied Data Science and Artificial Intelligence at Breda University of Applied Sciences, where I study on a scholarship with a 9.1/10 average.

Genmab appeals to me because its work stays close to both research and patients. Contributing to medicine in any form matters a lot to me, and working on data behind cancer treatment would be very motivating. I also like that this project has a clear practical aim: if historical controls can be reused safely, future studies can use fewer animals and still keep their statistical power.

The closest work I have done is a Kaggle competition on tracking cells in 3D microscopy of developing embryos, where my team of two finished in the top 14% of about 4,000 teams. The data came from only a few embryos, so we validated by leaving one embryo out. When hold-outs from embryos the model had already seen started ranking our changes wrongly, we built a cross-embryo check that agreed with the leaderboard. We also measured the noise floor of our scores and dropped settings whose effect changed sign from one embryo to another. This project asks a similar question in a more structured setting: how much of the variation comes from the animal and how much from the study, and whether a model still holds when one study is left out.

Statistical modelling, including Bayesian methods, was part of my first year, and in later projects we used statistical models as the baseline before moving on to machine learning. I would like to take this further into hierarchical models, which is a large part of why this project interests me. I have also worked with biological data before. I ranked 2nd of 77 in a Kaggle competition on plant root segmentation for NPEC, and in a five-person team I worked on a plant phenotyping platform where researcher feedback on flagged predictions drove retraining.

I enjoy working with people from other fields. In my current project with CEVA Logistics and bol, I work with several companies to get access to their data, help write business requirements documents, and make sure our proposals are clear to people without a technical background. I ask a lot of questions when I am new to a field, and I expect to do that often with the scientists in Non-Clinical Safety and Pre-clinical in vivo Pharmacology.

I work mainly in Python, with Pandas and NumPy for data work. I use Git for everything and keep my code tested and documented; the platform project ran an 85% coverage gate across 430 tests. Utrecht is a direct train ride from Tilburg, so three days a week on site works well for me.

I would be glad to talk about how I could contribute to the project. Thank you for your time and consideration.

Kind regards,

Oleksii Krasnoshtanov